

# Our Chat !

**An easy-to-install Enterprise LLM chat system using LibreChat with AWS Bedrock and LDAP/AD authentication.**

Why is this needed? Can't users just access AWS Bedrock directly? They might; however, in most organizations, the AWS Console is reserved for power users, as it typically takes some time until you are familiar with it (to put it mildly).

LibreChat, on the other hand, takes zero on-boarding time; users simply login with their enterprise credentials and use the system just like ChatGPT or Claude.ai. Another reason to like LibreChat is its superior user interface. It has gained much popularity and is often trending on GitHub.

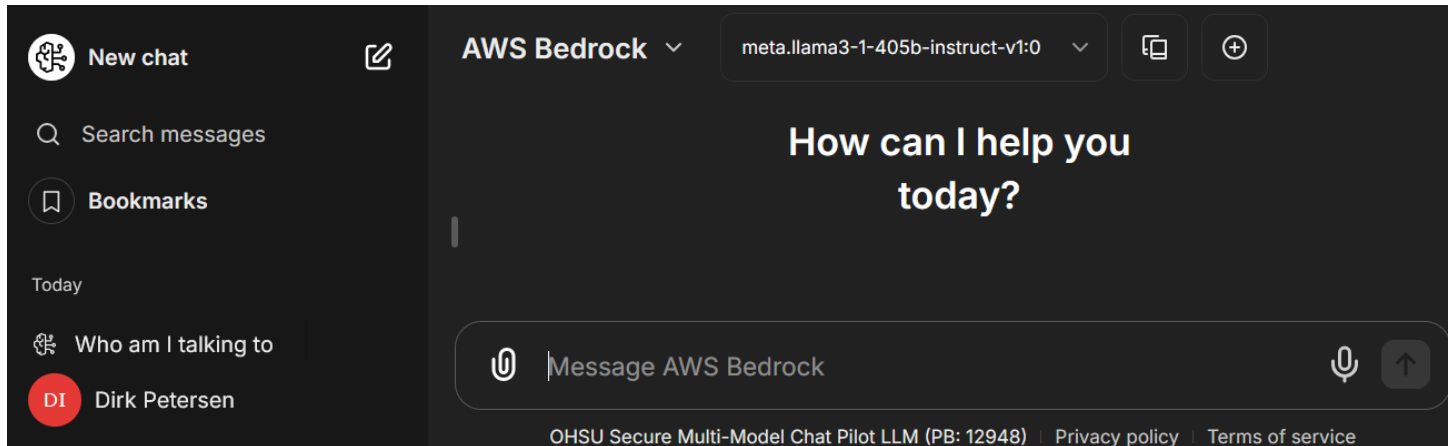


Figure 1: image

Other big benefits:

- No per-user cost. Even a few hundred ChatGPT users can cost an organization hundreds of thousands of dollars per year. With LibreChat, you can serve tens of thousands of users and only pay the API cost of what they actually use.
- No configuration needed to prevent the cloud LLM from learning with your data. All API use in Bedrock, as well as OpenAI, is excluded from learning usage.
- Presumably better security: OpenAI, Anthropic, and other startups have significantly smaller cybersecurity operations than the cloud Hyperscalers, such as AWS, Azure, & Google.
- Unified user interface: In addition to Bedrock, you can use OpenAI or Google API or even on-prem use cases, all within the same interface.

Table of Contents:

- Prerequisites
- Prepare Server
- Prepare install
- AWS connectivity
- SSL certificates
- LibreChat Configuration
  - .env
  - librechat.yml (optional)
  - nginx.conf (optional)
- Install LibreChat
- Budgeting
- API usage
- Purging of old chats
- Updating Librechat
- Disaster Recovery / Business Continuity

- On-premises use cases
- Longer term vision
- Troubleshooting
  - Get debug output
  - cannot create docker group
  - I don't have root permissions
  - Cleaning up docker

## Prerequisites

- Get a RHEL virtual server (this process was tested with RHEL 9.4) with at least 8GB RAM and 50GB free disk space, of which about half should be under /home. As this server might process sensitive data, ask for all security software (log forwarder, Antivirus/malware, intrusion prevention) to be preinstalled.
- This machine must be able to communicate with the `ldaps port 636` of your enterprise LDAP server (for example, Active Directory).
- An LDAP/AD security group that contains the users who are allowed to use the chat system. For now, we call this group `our-chat-users`.
- An SSL certificate, unless you use Let's Encrypt.
- AWS credentials (`AWS_ACCESS_KEY_ID`, `AWS_SECRET_ACCESS_KEY`) for an AWS service account (perhaps called `librechat` or `ochat`) that has no permissions except for the `AmazonBedrockFullAccess` policy attached to it.
- You don't require root access if your sysadmins can run the `prepare-server.sh` script for you, but they should allow you to switch to the `ochat` user, e.g., `sudo su - ochat`.

## Prepare Server

Run the `prepare-server.sh` script as root user to install docker and prepare the `ochat` user account. You can also start it as a normal user if you have requested the correct sudo config.

```
curl https://raw.githubusercontent.com/dirkpetersen/our-chat/refs/heads/main/prepare-server.sh?token=$(date +%s)
```

Now switch to the `ochat` user with `sudo su - ochat` and continue with Prepare install.

## Prepare install

After switching to the `ochat` user (`sudo su - ochat`), please clone the `our-chat` repository from GitHub and copy the `.env.ochat`, `librechat.yml` and `nginx.conf` files to the root of the home directory of the `ochat` user:

```
cd ~
git clone https://github.com/dirkpetersen/our-chat/
cp ~/our-chat/.env.ochat ~/.env
cp ~/our-chat/librechat.yaml ~/librechat.yaml
cp ~/our-chat/nginx.conf ~/nginx.conf
```

## AWS connectivity

Export the AWS credentials and region of the Bedrock service account (the one that has only the `AmazonBedrockFullAccess` policy attached) to environment variables:

```
export AWS_ACCESS_KEY_ID=abcdefghijklmnopqrstuvwxyz
export AWS_SECRET_ACCESS_KEY=abcdefghijklmnopqrstuvwxyz123456
export AWS_DEFAULT_REGION=us-west-2
```

Then test the connectivity to AWS Bedrock by running:

```
~/our-chat/tests/bedrock-test.py
```

You should see that your credentials have been written to the correct places under `~/aws` and Bedrock shows a list of available models and responds to a "Hello World" prompt:

```
Written to /home/librechat/.aws/credentials: AWS credentials for [default]
Written to /home/librechat/.aws/config: AWS region configuration for [default]
```

List of available models on AWS Bedrock:

```
amazon.titan-tg1-large
amazon.titan-embed-g1-text-02
.
.
mistral.mistral-large-2407-v1:0
```

Response to 'Hello, world': Hello! How can I assist you today? Feel free to ask me anything or let me know if

If you don't get that response or the script shows an error, go back to your AWS Administrator for troubleshooting your AWS credentials or permissions before you continue.

You can find more details about AWS in the AWS budget section below.

## SSL certificates

Here we cover the standard case of you receiving SSL certs from your enterprise team. In many cases your team will send you a PKCS12 archive which comes as a \*.pfx file along with an import password. After entering the import password, you need to set a new PEM pass phrase. At the end of this process you should have ~/our-chat.pem and ~/our-chat.pw at the root of the ochat home directory. (Please note: If you **leave a space at the beginning** of the echo "your-pem-passphrase" line, the pass phrase will not end up in your bash history).

```
openssl pkcs12 -in cert-from-IT.pfx -out ~/our-chat.pem
```

Enter Import Password:

Enter PEM pass phrase:

Verifying - Enter PEM pass phrase:

```
echo "your-pem-passphrase" > ~/our-chat.pw
chmod 600 ~/our-chat.*
```

## LibreChat Configuration

You will likely need to edit each of these config files at some point, but for now you should only edit the ~/.env file to enable LDAP authentication and to update a few security tokens. If you are in a Disaster Recovery / Business Continuity situation, please do not edit the 3 files below but restore them from your secure location.

### .env

Please find these settings in ~/.env:

```
vi ~/.env
```

```
LDAP_URL=ldaps://ldap.domain.edu:636
LDAP_USER_SEARCH_BASE=OU=User Accounts,dc=domain,dc=edu
LDAP_BIND_DN=CN=myserviceaccount,OU=Service Accounts,OU=User Accounts,DC=domain,DC=edu
LDAP_BIND_CREDENTIALS="ad-password"
LDAP_LOGIN_USES_USERNAME=true
LDAP_SEARCH_FILTER=(&(sAMAccountName={{username}})(memberOf=CN=our-chat-users,OU=Groups,DC=domain,DC=edu))
LDAP_FULL_NAME=displayName
```

These settings should be pretty self-explanatory. For example, LDAP\_LOGIN\_USES\_USERNAME means that you can log in with your username instead of typing your entire email address. The LDAP\_SEARCH\_FILTER is a bit of a workaround that

we use to restrict LibreChat to members of an AD/LDAP security group. The filter was not intended for authorization, and if a new user is not a member of that group, they will get a 401 error (AuthN) instead of 403 (AuthZ). This can be a bit confusing. On some LDAP systems, the LDAP\_BIND\_DN can be the email address (aka service principal) of the service account, e.g., `myserviceaccount@domain.edu`, but it is safest to use the DN (distinguished name).

After you have configured all these 7 settings, please use the LDAP test script `~/our-chat/tests/ldap-test.py <testuser>` to verify these settings. To check the LDAP\_SEARCH\_FILTER, you have to pass the username of a test user who is member of that security group (e.g. `our-chat-users`)

```
~/our-chat/tests/ldap-test.py peterdir
```

```
/ldap-test.py
```

```
Successfully read environment variables: ['LDAP_URL', 'LDAP_USER_SEARCH_BASE', 'LDAP_BIND_DN', 'LDAP_BIND_CRED
```

```
Connected as: CN=myserviceaccount,OU=Service Accounts,OU=User Accounts,DC=domain,DC=edu
```

```
Evaluating LDAP_SEARCH_FILTER with testuser peterdir:
```

```
Evaluating search filter: (&(sAMAccountName=peterdir)(memberOf=CN=our-chat-users,OU=Groups,DC=domain,DC=edu))
```

```
Found 1 matching entries:
```

```
DN: CN=peterdir,OU=User Accounts,DC=ohsum01,DC=ohsu,DC=edu
```

```
Attributes:
```

```
    displayName: Dirk Petersen
```

```
    memberOf: ['CN=unix_ochat,OU=Groups,DC=domain,DC=edu', .....
```

And as a final step we want to setup unique tokens for

CREDS\_KEY, CREDS\_IV, JWT\_SECRET, JWT\_REFRESH\_SECRET and MEILI\_MASTER\_KEY

Go to [https://www.librechat.ai/toolkit/creds\\_generator](https://www.librechat.ai/toolkit/creds_generator), generate keys and replace the existing tokens for these 5 keys in `~/env`

```
vi ~/.env
```

You can likely skip the next 2 config files for now, but might want to change them later

### librechat.yml (optional)

For example, use `vi ~/librechat.yml && cp ~/librechat.yml ~/LibreChat/librechat.yml` to change the terms of service, modify the site footer and change a few advanced bedrock settings, for example allowed AWS regions.

### nginx.conf (optional)

The only change `vi ~/nginx.conf && cp ~/nginx.conf ~/LibreChat/client/nginx.conf` likely requires, is setting the filenames for the SSL certificates for https, if you choose a different certificate name than `our-chat.pem` / `our-chat.pw`. You can set additional nginx headers to further enhance security.

```
ssl_certificate /etc/librechat/ssl/our-chat.pem;
ssl_certificate_key /etc/librechat/ssl/our-chat.pem;
ssl_password_file /etc/librechat/ssl/our-chat.pw;
```

## Install LibreChat

When all the prep work is done correctly, you should be able to run `install-librechat.sh` and have a running system in a few seconds:

```
~/our-chat/install-librechat.sh
```

Cloning into 'LibreChat'...

```

remote: Enumerating objects: 33792, done.
remote: Counting objects: 100% (5880/5880), done.
remote: Compressing objects: 100% (1247/1247), done.
remote: Total 33792 (delta 5103), reused 4786 (delta 4622), pack-reused 27912 (from 1)
Receiving objects: 100% (33792/33792), 43.72 MiB | 44.24 MiB/s, done.
Resolving deltas: 100% (24182/24182), done.
Copying /home/ochat/LibreChat/deploy-compose.yml to /home/ochat/LibreChat/deploy-compose-ourchat.yml
~/aws/src has been added to .bashrc
Copying /home/ochat/.env to /home/ochat/LibreChat/.env and expanding env vars
Copying /home/ochat/librechat.yaml to /home/ochat/LibreChat/librechat.yaml
Copying /home/ochat/nginx.conf to /home/ochat/LibreChat/client/nginx.conf
Generating DH parameters, 2048 bit long safe prime
.
.
[+] Running 8/8
    Network librechat_default      Created
    Volume "librechat_pgdata2"     Created
    Container chat-mongodb         Started
    Container chat-meilisearch     Started
    Container librechat-vectordb-1 Started
    Container librechat-rag_api-1  Started
    Container LibreChat-API        Started
    Container LibreChat-NGINX      Started

no crontab for ochat
Cron job added to run /home/ochat/purge_old_messages.py daily at 2:22 AM

stopping: docker compose -f /home/ochat/LibreChat/deploy-compose-ourchat.yml down
starting: docker compose -f /home/ochat/LibreChat/deploy-compose-ourchat.yml up -d

Now try to access your chat system, e.g. https://ourchat.domain.edu. If you encounter issues, please see the troubleshooting section below.

```

## Budgeting

Since LibreChat does not support cost control at this time, we need to rely on AWS. The service **AWS Budgets** is a good start, but we quickly realize that this product is for alerting only and does not actually stop services from being used if there is a budget overrun. AWS will point out that there is another service called **AWS Budget Actions**, but it seems those are mainly triggering **AWS Lambda** scripts. They also lack the flexibility we need, which raises the question: why not just run an hourly Lambda script to check our spend?

Let's assume our monthly budget is \$1000 (and not a penny more). If we disable the service after \$950 has been spent, this may happen on the first day of the month, and then nobody will be able to use the system for the rest of the month. We could also set the budget to \$32 per day, and with a maximum of 31 days in a month, we would never exceed \$992 / month. But that would mean people cannot really do big things, and we want to support innovation. We do not want someone to be constrained by the \$32 per day if only \$100 has been spent this month.

So let's see if we can accumulate the budget. For example, if the system has not been used for 3 days, we would set the budget of the 4th day to \$128 (\$32\*4). If someone came and used that entire amount on the 4th day, we would start fresh on the 5th day, and the budget would be \$32 per day again. In an extreme case, assume the system has not been used for the first 30 days in a month. Then someone could come along and use \$960 (\$32\*30) on the 31st day, and we would still be within budget.

How should this be implemented? Please see <https://github.com/dirkpetersen/our-chat/issues/1> for further discussion.

## API usage

LibreChat does not support API access (there is the RAG API feature, but that is a bit of a misnomer, as it actually requires the web ui). Instead, LibreChat is an aggregation point for many different APIs for users. Currently, LibreChat does not support complex cost control and all users of one server will access the same AWS account and there is only one single line item on the invoice and we cannot who was responsible for the cost. API users can automate their workflows which can lead to runaway costs. This means that API users should get their own AWS account with their own budget control and access the Bedrock API directly. `~/our-chat/tests/bedrock-test.py` may service as a resonable initial example that can list LLMs and ask a question.

## Purging of old chats

Especially in healthcare environments we want to make sure that sensitive data does not reside on systems any longer than necessary. As of October 2024 LibreChat does not have the ability to purge data, however OurChat by default has a purge script activated, that deletes any messages and files older than X days. (60 days by default)

## Updating Librechat

Update Librechat top the latest version and check the console output:

```
cd ~/LibreChat
docker compose -f ./deploy-compose-ourchat.yml down
git pull
docker compose -f ./deploy-compose-ourchat.yml pull
docker compose -f ./deploy-compose-ourchat.yml up
```

if everything works, we start LibreChat in Daemon mode:

```
docker compose -f ~/LibreChat/deploy-compose-ourchat.yml down
docker compose -f ~/LibreChat/deploy-compose-ourchat.yml up -d
```

should an upgrade fail, please follow these steps (Note: The chat history and uploaded files of all users will be lost)

```
docker compose -f ~/LibreChat/deploy-compose-ourchat.yml down
mv ~/LibreChat ~/LibreChat.fail.1
~/our-chat/install-librechat.sh
```

## Disaster Recovery / Business Continuity / Emergency Operation

As the system only stores data temporarily, disaster recovery and business continuity procedures are limited to backing up 5 configuraton files / certificates in a secure place:

```
~/ .env
~/librechat.yaml
~/nginx.conf
~/our-chat.pem
~/our-chat.pw
```

The install procedures at <https://github.com/dirkpetersen/our-chat> also serve as disaster recovery process as they are enabling you to restore service in less than 30 minutes, once you have the prerequisites met. (You may decide to have another VM already running to reduce the provisioning time).

While following the install procedures do not edit any of the 5 files but restore them from your secure backup location to `~/` (the root of the home direcotry of the service account) and copy them to the right locations in your LibreChat installation:

```
cp ~/ .env ~/LibreChat/
cp ~/librechat.yaml ~/LibreChat/
cp ~/nginx.conf ~/LibreChat/client/
cp ~/our-chat.pem ~/LibreChat/client/ssl/
```

```
cp ~/our-chat.pw ~/LibreChat/client/ssl/
```

## On-premises use cases

If you have use cases that are not permitted to use cloud services (e.g. patent issues), you can use on premises GPUs or even connect to your HPC cluster and use those resources all within the same LibreChat interface. Configure these custom endpoints in `.env` and in `librechat.yaml`

## Longer term vision

In the future we may extend our chat eco-system, and add apps with additional capabilities but at this time LibreChat seems to meet our needs.

## Troubleshooting

### Get debug output

If something is not working, the first step is to enable debugging and bringing up the docker containers in non-daemon mode so that they are printing all logs to the console. Setup debug output, open `~/env` in editor

```
vi ~/.env && cp ~/.env ~/LibreChat/
```

and set `DEBUG_LOGGING=true`

Open up a second terminal, shutdown the containers and bring them up again, this time (up without a `-d`)

```
docker compose -f ~/LibreChat/deploy-compose-ourchat.yml down
```

```
docker compose -f ~/LibreChat/deploy-compose-ourchat.yml up
```

Check for error messages. When you are done, execute the `down` and then the `up -d` command

### cannot create docker group

In some cases your docker installation may not create a `docker` group in `/etc/group`. This is often because your IT department may have created a global docker group in their IAM system in order to manage their systems centrally. In that case the `groupadd docker` command will fail. There are 2 workarounds for this:

1. You can try editing `/etc/group` directly and add a line `docker:x:986:ochat` to it. The downside of this approach is, that you then have 2 different docker groups with different `gidNumbers`. While it should work without issues, it can be confusing, especially if you are troubleshooting
2. You can create a new group (for example `ldocker`, e.g. for local docker), add the `ochat` user to it and ask docker to use it instead of the `docker` group. This requires these steps:

- edit `/etc/docker/daemon.json`

```
sudo echo -e '{\n  "group": "ldocker"\n}' >> /etc/docker/daemon.json
```

- run `systemctl edit docker.socket` and set the `SocketGroup=ldocker`

```
### Editing /etc/systemd/system/docker.socket.d/override.conf
```

```
### Anything between here and the comment below will become the new contents of the file
```

```
[Socket]
```

```
SocketGroup=ldocker
```

```
### Lines below this comment will be discarded
```

- restart the docker socket (not the docker daemon) and check that the socket has the correct group ownership

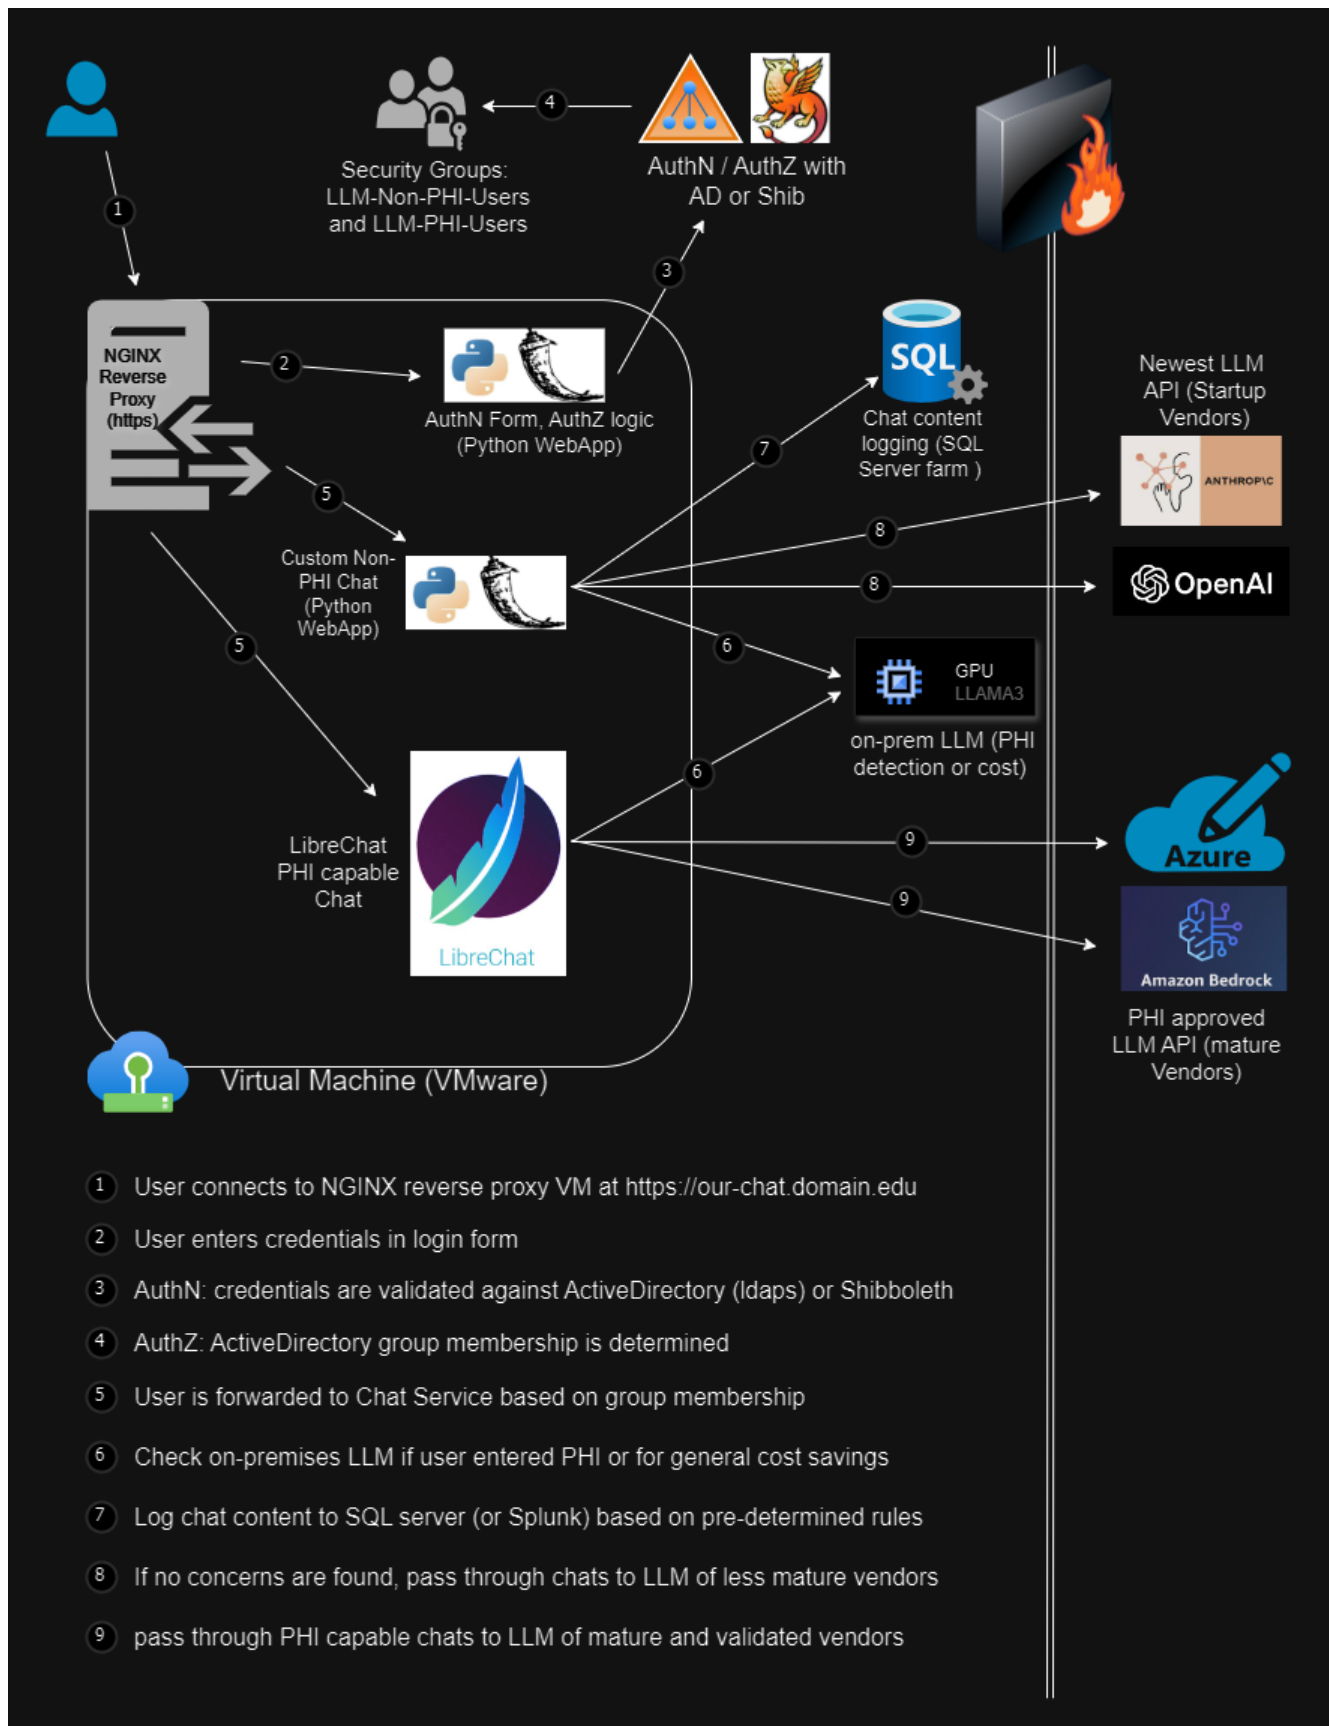


Figure 2: our-chat-dark



```
systemctl restart docker.socket
```

```
ls -l /var/run/docker.sock  
srw-rw----. 1 root ldocker 0 Oct  4 21:23 /var/run/docker.sock
```

If this is not working, please remove the docker packages using **dnf** and reinstall docker

### **I don't have root permissions**

If your IT infrastructure team cannot give you **root** access to the virtual server you requested, you may still be able to get access to management features via **sudo**. Ask your sysadmin to run **visudo** and paste in the config below. Change **yourusername** to your actual user name:

```
yourusername (ALL) NOPASSWD: /usr/bin/dnf, /usr/bin/systemctl, /usr/bin/loginctl enable-linger *, /usr/bin/doc  
yourusername (ALL) !/usr/bin/su -  
yourusername (ALL) !/usr/bin/su - root
```

optionally you could also get these, for example if you need another NGINX reverse proxy

```
yourusername (ALL) NOPASSWD: /usr/bin/rpm, /usr/bin/vi /etc/nginx/nginx.conf, /usr/bin/vi /etc/nginx/sites-ava
```

### **Cleaning up docker**

```
### cleanup as root  
sudo systemctl stop docker  
sudo rm -rf /var/lib/docker/*  
sudo rm -rf /var/lib/containerd/*  
sudo systemctl start docker
```